

Class 7: Machine Learning 1

Barry

Table of contents

| | |
|--|----|
| Clustering | 1 |
| K-means | 4 |
| Hierarchical Clustering | 7 |
| Principal Component Analysis (PCA) | 9 |
| Data import | 9 |
| PCA to the rescue | 12 |
| Interperiting PCA Results | 12 |


Today we will explore unsupervised machine learning methods starting with clustering and dimensionality reduction.

Clustering

To start let's make up some data to cluster where we know what the answer should be. The `rnorm()` function will help us here.

```
hist( rnorm(10000, mean=3) )
```

Histogram of rnorm(10000, mean = 3)



Return 30 numbers centred on -3

```
tmp <- c( rnorm(30, mean=-3),  
          rnorm(30, mean=+3) )  
  
x <- cbind(x=tmp, y=rev(tmp))  
  
x
```


| | x | y |
|-------|------------|-----------|
| [1,] | -4.0052498 | 4.3175912 |
| [2,] | -2.5365722 | 1.8276559 |
| [3,] | -3.5911897 | 4.0986464 |
| [4,] | -2.4508472 | 2.5887270 |
| [5,] | -4.1325932 | 2.0106003 |
| [6,] | -3.7806618 | 1.6829283 |
| [7,] | -3.3607000 | 2.0327000 |
| [8,] | -2.6114337 | 0.7115766 |
| [9,] | -2.0864165 | 1.1543562 |
| [10,] | -5.5458617 | 4.8342991 |
| [11,] | -1.1335161 | 3.1381491 |
| [12,] | -2.2518117 | 2.8710214 |
| [13,] | -2.1482412 | 2.9451070 |

| | | |
|-------|------------|------------|
| [14,] | -2.2651313 | 1.6493587 |
| [15,] | -2.5145047 | 1.7154397 |
| [16,] | -1.8756006 | 3.8184801 |
| [17,] | -3.7421402 | 3.1095511 |
| [18,] | -3.3345238 | 3.2019636 |
| [19,] | -2.2717764 | 2.9436002 |
| [20,] | -4.0122721 | 2.9230563 |
| [21,] | -2.3392715 | 3.8003348 |
| [22,] | -3.3265262 | 2.2767546 |
| [23,] | -2.9726308 | 2.5225962 |
| [24,] | -1.9712435 | 4.5665173 |
| [25,] | -2.0126858 | 4.1387376 |
| [26,] | -4.3339338 | 2.9478143 |
| [27,] | -3.3677454 | 3.0854448 |
| [28,] | -2.2030091 | 2.0877410 |
| [29,] | -5.5590850 | 2.2508092 |
| [30,] | -2.6464799 | 3.4467805 |
| [31,] | 3.4467805 | -2.6464799 |
| [32,] | 2.2508092 | -5.5590850 |
| [33,] | 2.0877410 | -2.2030091 |
| [34,] | 3.0854448 | -3.3677454 |
| [35,] | 2.9478143 | -4.3339338 |
| [36,] | 4.1387376 | -2.0126858 |
| [37,] | 4.5665173 | -1.9712435 |
| [38,] | 2.5225962 | -2.9726308 |
| [39,] | 2.2767546 | -3.3265262 |
| [40,] | 3.8003348 | -2.3392715 |
| [41,] | 2.9230563 | -4.0122721 |
| [42,] | 2.9436002 | -2.2717764 |
| [43,] | 3.2019636 | -3.3345238 |
| [44,] | 3.1095511 | -3.7421402 |
| [45,] | 3.8184801 | -1.8756006 |
| [46,] | 1.7154397 | -2.5145047 |
| [47,] | 1.6493587 | -2.2651313 |
| [48,] | 2.9451070 | -2.1482412 |
| [49,] | 2.8710214 | -2.2518117 |
| [50,] | 3.1381491 | -1.1335161 |
| [51,] | 4.8342991 | -5.5458617 |
| [52,] | 1.1543562 | -2.0864165 |
| [53,] | 0.7115766 | -2.6114337 |
| [54,] | 2.0327000 | -3.3607000 |
| [55,] | 1.6829283 | -3.7806618 |
| [56,] | 2.0106003 | -4.1325932 |

```
[57,] 2.5887270 -2.4508472
[58,] 4.0986464 -3.5911897
[59,] 1.8276559 -2.5365722
[60,] 4.3175912 -4.0052498
```

Make a plot of x

```
plot(x)
```



K-means

The main function in “base” R for K-means clustering is called `kmeans()`:

```
km <- kmeans(x, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:


| | x | y |
|---|-----------|----------|
| 1 | -3.012788 | 2.823278 |


```
km$centers
```

| | x | y |
|---|-----------|-----------|
| 1 | -3.012788 | 2.823278 |
| 2 | 2.823278 | -3.012788 |


Q. Make a plot of our `kmeans()` results showing cluster assignment using different colors for each cluster/group of points and cluster centers in blue.

```
plot(x, col=km$cluster )  
points(km$centers, col="blue", pch=15, cex=2)
```



Q. Run `kmeans()` again on `x` and this cluster into 4 groups/clusters and plot the same result figure as above.

```
km4 <- kmeans(x, centers = 4)  
plot(x, col=km4$cluster )  
points(km4$centers, col="blue", pch=15, cex=2)
```



key-point: K-means clustering is super popular but can be misused. One big limitation is that it can impose a clustering pattern on your data even if clear natural grouping doesn't exist - i.e. it does what you tell it to do in terms of **centers**.

Hierarchical Clustering

The main function in “base” R for hierarchical clustering is called `hclust()`.

You can't just pass our dataset as is into `hclust()` you must give “distance matrix” as input. We can get this from the `dist()` function in R.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

The results of `hclust()` don't have a useful `print()` method but do have a special `plot()` method.

```
plot(hc)
abline(h=8, col="red")
```



To get our main cluster assignment (membership vector) we need to “cut” the tree at the big goal posts...

```
grps <- cutree(hc, h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
```


```
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
table(grps)
```

```
grps
  1  2
30 30
```



```
plot(x, col=grps)
```



Hierarchical clustering is distinct in that the dendrogram (tree figure) can reveal the potential grouping in your data (unlike K-means)

Principal Component Analysis (PCA)

PCA is a common and highly useful dimensionality reduction technique used in many fields - particularly bioinformatics.

Here we will analyze some data from the UK on food consumption.

Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)

head(x)
```

| | X | England | Wales | Scotland | N.Ireland |
|---|---------------|---------|-------|----------|-----------|
| 1 | Cheese | 105 | 103 | 103 | 66 |
| 2 | Carcass_meat | 245 | 227 | 242 | 267 |
| 3 | Other_meat | 685 | 803 | 750 | 586 |
| 4 | Fish | 147 | 160 | 122 | 93 |
| 5 | Fats_and_oils | 193 | 235 | 184 | 209 |
| 6 | Sugars | 156 | 175 | 147 | 139 |


```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

| | England | Wales | Scotland | N.Ireland |
|---------------|---------|-------|----------|-----------|
| Cheese | 105 | 103 | 103 | 66 |
| Carcass_meat | 245 | 227 | 242 | 267 |
| Other_meat | 685 | 803 | 750 | 586 |
| Fish | 147 | 160 | 122 | 93 |
| Fats_and_oils | 193 | 235 | 184 | 209 |
| Sugars | 156 | 175 | 147 | 139 |


```
x <- read.csv(url, row.names = 1)
head(x)
```

| | England | Wales | Scotland | N.Ireland |
|---------------|---------|-------|----------|-----------|
| Cheese | 105 | 103 | 103 | 66 |
| Carcass_meat | 245 | 227 | 242 | 267 |
| Other_meat | 685 | 803 | 750 | 586 |
| Fish | 147 | 160 | 122 | 93 |
| Fats_and_oils | 193 | 235 | 184 | 209 |
| Sugars | 156 | 175 | 147 | 139 |

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



One conventional plot that can be useful is called a “paris” plot.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



PCA to the rescue

The main function in base R for PCA is called `prcomp()`.

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 |
|------------------------|----------|----------|----------|-----------|
| Standard deviation | 324.1502 | 212.7478 | 73.87622 | 2.921e-14 |
| Proportion of Variance | 0.6744 | 0.2905 | 0.03503 | 0.000e+00 |
| Cumulative Proportion | 0.6744 | 0.9650 | 1.00000 | 1.000e+00 |

Interpreting PCA Results

The `prcomp()` function returns a list object of our results with five attributes/components

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

The two main “results” in here are `pca$x` and `pca$rotation`. The first of these (`pca$x`) contains the scores of the data on the new PC axis - we use these to make our “PCA plot”.

```
pca$x
```


| | PC1 | PC2 | PC3 | PC4 |
|-----------|------------|-------------|------------|---------------|
| England | -144.99315 | -2.532999 | 105.768945 | -9.152022e-15 |
| Wales | -240.52915 | -224.646925 | -56.475555 | 5.560040e-13 |
| Scotland | -91.86934 | 286.081786 | -44.415495 | -6.638419e-13 |
| N.Ireland | 477.39164 | -58.901862 | -4.877895 | 1.329771e-13 |

```
library(ggplot2)
```

```
library(ggrepel)
```

```
# Make a plot of pca$x with PC1 vs PC2
```

```
ggplot(pca$x) +  
  aes(PC1, PC2, label=rownames(pca$x)) +  
  geom_point() +  
  geom_text_repel()
```



The second major result is contained in the `pca$rotation` object or component. Let's plot this to see what PCA is picking up...

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```

